Interpersonal Coordination in Perception and Memory in an Online Experiment: Using Networked Crowdsourcing for Experiments on Human Social Interaction

Alexandra Paxton
paxton.alexandra@gmail.com
Institute of Cognitive and Brain Sciences
Berkeley Institute for Data Science
University of California, Berkeley

Thomas J. H. Morgan
thomas.j.h.morgan@asu.edu
School of Human Evolution and Social Change
Arizona State University

Jordan W. Suchow
suchow@berkeley.edu
Social Science Matrix
University of California, Berkeley

Thomas L. Griffiths
tom_griffiths@berkeley.edu
Department of Psychology
University of California, Berkeley

Abstract

Recent advances in crowdsourcing have helped many cognitive scientists reach out beyond traditional undergraduate subject pools to run a range of experimental paradigms with a wider audience. To date, however, many of these opportunities for online experiments on crowdsourcing platforms have been closed to researchers interested in capturing the dynamics of human social interaction. We argue that an important next step for increasing the adoption and utility of online experiments will lie in using *networked crowdsourcing*—moving beyond providing individual participants separate tasks to support more complex interactive or interdependent configurations. Networked crowdsourcing allows researchers to capture real-time and transmission-chain interaction between participants to study social cognition and behavior. Here, we use networked crowdsourcing to move the study of real-time interpersonal coordination from the lab and onto Amazon Mechanical Turk, examining how people grow similar over time in their perception and memory.

Keywords: interpersonal coordination; networked crowdsourcing; human communication; online experiments; social interaction

Feedback from Tom

decrease the "network" stuff – focus on online expeimrnts

Introduction

Research on the phenomenon of *interpersonal coordination* focuses on the subtle ways in which our interactions with others directly affect our own behaviors, feelings, and thoughts. Interest has surged over the last several decades in understanding how contact with others shapes our cognition and behavior. Much of the research on coordination (also known as interactive alignment, interpersonal synchrony, mimicry, and more; see Paxton, Dale, & Richardson, 2016) focuses on how we become more similar over time, especially during task-oriented or friendly contexts.

A growing perspective in this area has taken inspiration from dynamical systems theory, conceptualizing interaction as a complex adaptive system from which coordination arises as an emergent phenomenon according to contextual pressures (Riley, Richardson, Shockley, & Ramenzoni, 2011). A fundamental principle of this dynamical systems perspective holds that coordination should not be static across contexts nor over time. Exploring how new contexts and contextual demands—like interpersonal conflict (Paxton & Dale, 2013), friendly competition (Tschacher, Rees, & Ramseyer, 2014), or specialized task demands (Ramenzoni, Riley, Shockley, & Baker, 2012)—has become a central part of this perspective, laying out under what conditions coordination disappears, increases, or demonstrates complementary rather than identical patterns (Fusaroli et al., 2012).

Research is similarly interested in comparing how coordination changes across different behavioral or cognitive systems. Under the dynamical systems perspective, a context's unique pressures, the resulting coordination dynamics, and their impact on the interaction may differ—for example, demonstrating that individuals tend to become more similar over time across a variety of metrics (Louwerse, Dale, Bard, & Jeuniaux, 2012) but that specific kinds of coordination can help or hurt outcomes (Fusaroli et al., 2012) during task-related interaction.

Broadly, during tasks that are neutral (Shockley, Santana, & Fowler, 2003), cooperative (Louwerse et al., 2012), or non-confrontationally competitive (e.g., competitive games; Tschacher et al., 2014), previous work broadly suggests that individuals' behavior and cognition become more similar over time. A range of behavioral signals—both high-level (e.g., gesture; Louwerse et al., 2012) and low-level (e.g., postural sway; Shockley et al., 2003)—become synchronized, even when the interacting individuals are unable to see one another (Shockley et al., 2003) or are separated in time (D. C. Richardson & Dale, 2005).

The systematic testing of coordination across a variety of interaction contexts is vital to charting its dynamical landscape. This methodical exploration of different factors will eventually enable us to identify control parameters and key factors of initial conditions that shape how coordination emerges and how it impacts interaction outcomes. Doing so, however, requires researchers on interpersonal coordination to expand our view of experimental paradigms: Even as we continue to embrace more complex naturalistic interactions (e.g., Tschacher et al., 2014), we must continue developing experimental methods for analyzing "minimally interactive contexts" (Hale, Pan, & Hamilton, 2015)—that is, situations

in which our interactions with others are limited in behavioral channel, scope, or time—to fully map the interaction space.

In the search for both fully interactive and minimally interactive paradigms, we believe that online experiment platforms and crowdsourcing can serve as powerful tools in experiment development and deployment. By connecting people digitally, researchers are able to more fully control the experimental experience, from deciding how much social information about their partner(s) will be available to establishing what communication channel(s) can be utilized to crafting interactive studies for groups beyond the dyad. We propose that researchers interested in social cognition and behavior take advantage of the rapid advances in crowdsourcing platforms to develop *networked crowdsourcing* experiment models.

Networked Crowdsourcing for Social Experiments

Recent developments in how data can be collected and analyzed are transforming cognitive science. This is reflected in an increased interest in big data and naturally occurring datasets (Goldstone & Lupyan, 2016), such as social media activity and video game logs, which hold the promise of capturing behavior in the wild and providing a testing ground for key scientific theories (Paxton & Griffiths, 2017). While these data can provide a window into observational data about human behavior at a massive scale, technological advances are quickly expanding to accommodate new *experimental* paradigms as well. [**From AP: Not sure this paragraph is useful anymore**]

Crowdsourcing platforms like Amazon Mechanical Turk (`http://www.mturk.com`) have been extensively used as a means to collect data with relatively simple but robust experimental paradigms, like surveys (Buhrmester, Kwang, & Gosling, 2011) and mouse-tracking (Freeman, Dale, & Farmer, 2011). At first, work in this domain required researchers to use established survey creation tools, which were quick to do but constrained experimental designs, or to program bespoke experiments, which is more open ended but far more time-consuming. More recently, cognitive scientists have worked to create solutions to support the efficient creation of a wider range of experiments (e.g., Gureckis et al., 2016). As the community around online psychology experiments has grown, it has done so with the intent to broaden its reach (especially to researchers with less programming experience) and to continue to provide more powerful experimental tools.

To date, many of these experiments have focused on individuals, making it difficult for researchers to study social processes through online experiments. We believe that the next step in online experimentation, then, is to move to *networked crowdsourcing*, creating interactive or interdependent experimental paradigms that construct interacting networks of people to understand social processes and phenomena. In doing so, networked crowdsourcing can provide researchers interested in social behavior the opportunity to expand their experimental capabilities beyond the lab while not compromising on the richness and complexity of true interactive contexts.

Online experimental platforms have been capable of serving individual paradigms for the past several years, but what marks this idea as different from these paradigms is our focus on *networks*. Instead of handling participants as individuals, networked crowsourcing allows researchers to connect participants with one another—whether through direct, real-time interaction or through sequential transmission chains—to directly manipulate social dynamics online in the same way that is possible within the lab. As such, networked crowdsourcing is uniquely positioned to support experimental research into human social behavior at scale.

A recurring concern for using online experiments lies in its participant population. Like all convenience samples, there can be questions about the degree to which the participants reflect the broader population dynamics—including the use of undergraduate students at Western universities as participants in return for course credit, who often do not reflect global demographics (Henrich, Heine, & Norenzayan, 2010). Considerations of sampling and population representativeness are vital for any study, and researchers should carefully consider their sampling choices at the outset of their work. For those interested in using online participants (especially from Amazon Mechanical Turk), recent surveys suggest that U.S.-based MTurk workers are more diverse in a variety of ways than typical college students but not entirely reflective of the general U.S. population (e.g., Buhrmester et al., 2011; Paolacci & Chandler, 2014).

The Present Study

The current study focuses on understanding how interacting individuals become entrained in perception and memory over time, becoming a sort of "line estimation system"—just as two people in an lab experiment become a "tangram recognition system" (Dale, Richardson, & Kirkham, 2011). To do that, we isolate the dyad to a minimally interactive context, allowing participants to engage with one another solely by communicating line lengths

Method

All research activities were completed in compliance with oversight from Committee for the Protection of Human Subjects at the University of California, Berkeley.

Participants

Participants (n = 148) were individually recruited from Amazon Mechanical Turk to participate as dyads (n = 74). Participants were paired with one another according to the order in which they began the experiment. All participants were over 18 years of age and fluent English speakers (self-reported); participation was restricted to only recruit from participants located within the U.S. with a 95% HIT approval rate.[1]

[1] A measure of MTurk worker quality, capturing how often their work is rejected by a requester. A 95% HIT approval rate means that only 5% of all of their submitted HITs have been rejected.

The experiment lasted an average of 11.69 minutes (range: 7.98—21.34 minutes). All participants were paid $1.33 as base pay for finishing the experiment and earned a bonus of up to $2 for the entire experiment based on mean accuracy over all trials (mean = $1.80; range: $0.00—$1.95). Participants were not informed about the amount of performance-based bonus that they earned during the experiment.

Procedure

All data collection procedures were run on Amazon Mechanical Turk (`http://mturk.com`) using the experiment platform Dallinger (v3.4.1; `http://github.com/dallinger/Dallinger`). Code for the experiment is available on GitHub (`http://github.com/thomasmorgan/joint-estimation-game`), and the resulting experiment data are available on the OSF repository for the project (`https://osf.io/8fu7x/`).

Each participant was individually recruited on Amazon Mechanical Turk to play a "Line Estimation Memory Game" (advertisement: "Test your memory skills!"). Upon completing informed consent, participants were told that they would be playing a game in which they would be required to remember and re-create line lengths. Participants were informed that they would be complete their training trials individually and would then begin playing with a partner. Participants were given no information about their partner other than the guess that their partner made; no information about the partner's identity was shared.

In each trial, participants were shown 3 red lines, each of a different length (see figure; **NB**: add figure), and were asked to remember all three of them.[2] The 3 stimulus lines were displayed for 2 seconds then removed, providing participants with a blank screen for 0.5 seconds. Participants were then told which line to re-create (#1, #2, or #3) and were then given 1 second to submit their guess at how long the target line had been. To do so, participants were given a blank box and used their cursor to fill in the box with a blue line. All lines were presented within bounded boxes of 500 pixels (wide) by 25 pixels (high).

During training, participants were then shown the correct length of the target line (as a grey bar above their own guess) for 2 seconds. This was accompanied by a message telling the participant that they had guessed correctly ("Your guess was correct!") or incorrectly ("Your guess was incorrect") or that they had not submitted a guess within the 1-second time limit ("You didn't respond in time").

During testing, participants' stimulus viewing, waiting, and recreation times remained the same as during testing, but they no longer received information about whether their guess was correct or incorrect. Instead, after both participants had submitted their first guess, participants were shown their guess (in blue) above their partner's guess (in green). Both participants were then asked whether they wanted to change their own guess or to keep the guess they had submitted. If either participant in the dyad indicated that they wanted to change their guess, that participant was then allowed to change their guess (again with a 1-second time limit) *while* still being able to view their partner's guess. Participants who chose to keep their previously submitted guess was informed that their partner chose to submit a new guess and waited for the other participant to finish. At that point, participants were again allowed to change or keep their guess. This process continued until both participants chose to keep their guess.

Participants were informed that their final accuracy would only be calculated for their final guess. However, because they had no means to communicate with their partner about whether each would be accepting or changing their guesses, each participant could not have known whether their decision to keep the guess would have been their final guess for the trial.

For clarity, we will refer to each new stimulus set as a *trial* and to each submitted line length estimate within each trial as a *guess*. This means that some participants may have submitted multiple guesses per trial. The last submitted estimate—the one by which trial-level accuracy is calculated—will be referred to as the *final guess*.

All dyads completed 10 training trials (alone) and 15 test trials (with their partner). All training and test stimuli were randomly generated for each dyad, but both participants within the dyad were given the same stimuli. After participating, each individual participant was asked to complete a series of questionnaires about the game on a series of 1-10 Likert-style scales, including the perceived difficulty of the task, how engaged they were in the task, and questions about their own and their parner's cooperativeness and trustworthiness.

Analyses

Similarity To measure how participants' perceptual and memory systems became more similar over time, we calculated the cross-correlation coefficient of participants' guess errors across trials (Paxton & Dale, 2013), within a window of +/-5 guesses. Although cross-correlation produces information about leading and following behavior, we have no *a priori* expectations about which of the two participants would emerge as a leader (given they have no information about their partner nor any assigned roles), so our first-pass analyses ignore any directionality by averaging across the each incremental lags (i.e., leading/following in both participants' directions).

Accuracy Accuracy was measured as a ratio relative to the total possible error on a given target stimulus trial. That is, rather than taking a given guess's error relative to the total line length, error was calculated as the maximum *possible* error. For example, if the target stimulus was 60 units long, the maximum possible error for that trial would be 60, and the

[2] A pilot version of this study showed that participants performed at ceiling when given only 1 line to remember and recreate. The additional 2 lines were added to strictly increase the memory load, as opposed to adding difficulty in other ways (e.g., creating a moving stimulus).

participant's error would be calculated relative to that maximum possible error.

Results

All analyses were performed in R (R Core Team, 2016). Linear mixed-effects models were performed using the `lme4` package (Bates, Mchler, Bolker, & Walker, 2015) using the maximal random slope structure for each random intercept to achieve model convergence (Barr, Levy, Scheepers, & Tily, 2013). All main and interaction terms were centered and standardized prior to entry in the model, allowing the model estimates to be interpretable as effect sizes (Keith, 2005).

Similarity in ratings over time

Our first model tested whether individuals' ratings became more similar over time while accounting for training improvement. We constructed a linear mixed-effects model predicting cross-correlation coefficients of normalized error with absolute lag and training improvement, using maximal random effects structures for experiment and ratings of difficulty after the experiment. As predicted, we find that dyads were significantly and strongly coupled in their error ratings (β = -0.43, $p <$ 0.0001), with no effect of training improvement (β = `r. round(ccf_main_model_readable$Estimate[training_row],2)`, $p < 0.36$).

However, during the experiment, both participants are also learning to play the game better over time, in addition to influencing one another. To ensure that the similarity observed in our first model was not simply an artifact of both participants improving individually, we next constructed a linear mixed-effects model to test whether the relation between participants' (1) adaptation to their partner's perceptual estimation and memory and (2) own performance changed over time. We calculated the former using normalized error ("true error") and the latter as the "error" between the participants' guesses ("partner error"). We built a linear mixed-effects model predicting the difference between "true error" and "partner error" with trial number and participants' training improvement, including maximal random effects structure for participant and difficulty. We found a trend toward a significant effect in the expected direction, with the difference decreasing slightly but not statistically significantly over time (β = 0.04, $p < 0.092$). . Again, we found no effect of training improvement (β = `r. round(checking_rates_model_readable$Estimate[training_row],2)`, $p < 0.42$).

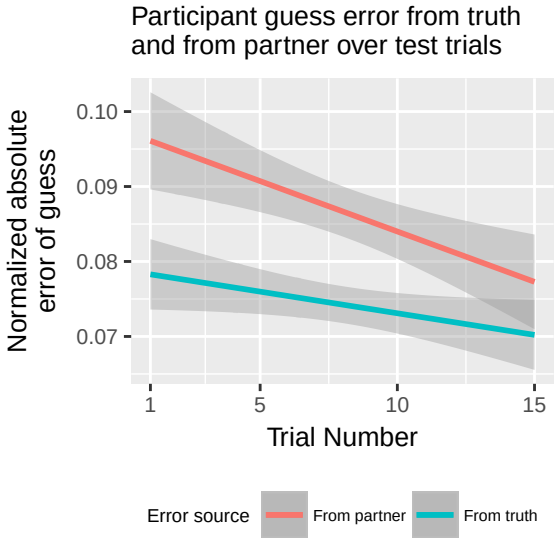

Figure 1: One column image.

Discussion

Future Directions

Injecting social dynamics in experimentally and/or as much as is desired

Conclusion

Acknowledgements

Thanks go to the Dallinger development team for their assistance in debugging the experiment.

This work was funded in part by the **DALLINGER GRANT INFO GOES HERE**.

References

Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure in mixed-effects models: Keep it maximal. *Journal of Memory and Language*, *63*(3), 255–278.

Bates, D., Mchler, M., Bolker, B., & Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, *67*(1), 1–48. `http://doi.org/10.18637/jss.v067.i01`

Buhrmester, M., Kwang, T., & Gosling, S. D. (2011). Amazon's Mechanical Turk: A new source of inexpensive, yet high-quality, data?, *6*, 3–5. `http://doi.org/10.2307/41613414`

Dale, R., Richardson, D. C., & Kirkham, N. K. (2011). *How two people become a tangram recognition system*. Aarhus University, Aarhus, Denmark.

Freeman, J. B., Dale, R., & Farmer, T. A. (2011). Hand in motion reveals mind in motion. *Frontiers in Psychology*, *2*, 59. `http://doi.org/10.3389/fpsyg.2011.00059`

Fusaroli, R., Bahrami, B., Olsen, K., Roepstorff, A., Rees, G., Frith, C., & Tyln, K. (2012). Coming to terms: Quanti-

fying the benefits of linguistic coordination. *Psychological Science*, *23*(8), 931–939.

Goldstone, R. L., & Lupyan, G. (2016). Discovering psychological principles by mining naturally occurring data sets. *Topics in Cognitive Science*, *8*, 548–568. http://doi.org/10.1111/tops.12212

Gureckis, T. M., Martin, J., McDonnell, J., Rich, A. S., Markant, D., Coenen, A., ... Chan, P. (2016). PsiTurk: An open-source framework for conducting replicable behavioral experiments online. *Behavior Research Methods*, *48*(3), 829–842.

Hale, J., Pan, X., & Hamilton, A. F. de C. (2015). Using interactive virtual characters in social neuroscience. In *Virtual reality (vr), 2015 ieee* (pp. 189–190). IEEE.

Henrich, J., Heine, S., & Norenzayan, A. (2010). Most people are not WEIRD. *Nature*, *466*(7302), 29.

Keith, T. (2005). *Multiple regression and beyond.* Boston, MA: Pearson.

Louwerse, M. M., Dale, R., Bard, E. G., & Jeuniaux, P. (2012). Behavior matching in multimodal communication is synchronized. *Cognitive Science*, *36*(8), 1404–1426.

Paolacci, G., & Chandler, J. (2014). Inside the Turk: Understanding Mechanical Turk as a participant pool. *Current Directions in Psychological Science*, *23*(3), 184–188. http://doi.org/10.1177/0963721414531598

Paxton, A., & Dale, R. (2013). Argument disrupts interpersonal synchrony. *Quarterly Journal of Experimental Psychology*, *66*(11). http://doi.org/10.1080/17470218.2013.853089

Paxton, A., & Griffiths, T. (2017). Finding the traces of behavioral and cognitive processes in big data and naturally occurring datasets. *Behavior Research Methods*, *49*(5), 1630–1638. http://doi.org/10.3758/s13428-017-0874-x

Paxton, A., Dale, R., & Richardson, D. (2016). Social coordination of verbal and nonverbal behaviors. In P. Passos, K. Davids, & C. J. Yi (Eds.), *Interpersonal coordination and performance in social systems* (pp. 259–273). Abington, Oxon; New York, NY: Routledge.

R Core Team. (2016). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from https://www.R-project.org/

Ramenzoni, V. C., Riley, M. A., Shockley, K., & Baker, A. A. (2012). Interpersonal and intrapersonal coordinative modes for joint and single task performance. *Human Movement Science*, *31*(5), 1253–1267.

Richardson, D. C., & Dale, R. (2005). Looking to understand: The coupling between speakers' and listeners' eye movements and its relationship to discourse comprehension. *Cognitive Science*, *29*(6), 1045–1060.

Riley, M. A., Richardson, M., Shockley, K., & Ramenzoni, V. C. (2011). Interpersonal synergies. *Frontiers in Psychology*, *2*, 38.

Shockley, K., Santana, M.-V., & Fowler, C. A. (2003). Mutual interpersonal postural constraints are involved in cooperative conversation. *Journal of Experimental Psychology: Human Perception and Performance*, *29*(2), 326.

Tschacher, W., Rees, G. M., & Ramseyer, F. (2014). Nonverbal synchrony and affect in dyadic interactions. *Frontiers in Psychology*, *5*, 1323.